

Deep Learning for NLP 2022

Homework 4

Due on June 13, 2021 at 23:59



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Version of May 31, 2022

Overall: 10 points.

This homework involves training more complex neural networks which can take some time. **Plan accordingly for the extended training time!**

Please note that this task requires a different conda environment with an older version of Tensorflow and Keras. **Setup a new conda environment with the requirements file in the homework archive and use it for this task only.** We apologize for the inconvenience.

Submission Guidelines

This homework comes with a zip archive containing the data and code templates. Write your code only in the `rnn.py` file. Then upload a zip file with the following folder structure:

```
hw4_group###.zip
├── hw4.pdf
└── rnn.py
```

- `hw4.pdf` should answer the questions **and** report your results.
- **Do NOT** include the subfolders in your submission (especially not the embeddings and your saved model). You can expect the embeddings to be available at `./embeddings/glove.6B.50d.txt`, the data in the `./data/` folder, and the utils at `./util/utils.py` (where “.” is the root folder of your zip).
- `###` is your group number (e.g., the file name could be `hw4_group2.zip`).
- Please make sure that your code prints **all** relevant results for **all** subtasks when the tutors run `rnn.py`.^a

Make sure that the results you report are reproducible and that your code is runnable in the **alternate conda environment**. If you are aware that your network never stops training, please be honest and add a short statement saying so. Thank you!

^aIf you think that it makes your code more readable/structured, you can also submit other `.py` files whose functions are called in `rnn.py`. But the tutors will only execute `rnn.py`.

1 Sequence Tagging with RNNs

(10P)

In this task, you will implement LSTM and Bi-LSTM architectures with Keras to perform part-of-speech tagging (a sequence tagging task).

Several Keras code snippets are provided for this home exercise. There is not much code to write in this exercise, but training RNNs is a computationally intensive task! **Plan accordingly for the extended training time.**

Data We use a subset of the data from the CoNLL-2003 shared task on Named Entity Recognition (provided in the zip). It is pre-partitioned into a training, development and test set.

The dataset consists of pre-tokenized sentences where every token is annotated with a part-of-speech tag, a syntactic chunk tag and a named entity tag. In this home exercise, we only use the part-of-speech tag.

1.1 Download Pretrained Embeddings

(0P)

Download the pretrained, uncased GloVe embeddings with 6B tokens (`glove.6B.zip`) from Stanford: <https://nlp.stanford.edu/projects/glove/>. For performance reasons, we will only use the 50-dimensional embeddings (`glove.6B.50d.txt`).

1.2 LSTM and Bi-LSTM Model

(2P)

Complete the network architecture by adding a respective LSTM and Bi-LSTM layer in the `build_model` method. Since the task is sequence tagging, your LSTM/Bi-LSTM should output a sequence. After the LSTM/Bi-LSTM, add a dense layer with n units to every sequence output, where n is the number of classes in the sequence tagging task plus 1.

Apply dropout before and after the LSTM layer.

Print the model summary in your code. **Copy-paste** or screenshot and paste the model summaries of both the LSTM and the Bi-LSTM into your PDF.

Hint:

- <https://keras.io/layers/recurrent/>

1.3 Intermediate Results

(2P)

Set the batch size to 10, dropout to 0.5, the number of hidden units to 100, and train your LSTM and Bi-LSTM on the training dataset for 10 epochs. **Print** and **report** (in your PDF) the final `val_categorical_accuracy` on the development set. **Print** and **report** the macro F1 score you achieve on the test set at this point in time.

1.4 F1 Model Checkpointer

(3P)

- The current implementation stores the model which produced the best `val_categorical_accuracy` on the development set. State in one sentence why this is a problematic approach and why a model checkpointer relying on macro F1 would be better suited. (1P)
- Complete the `on_epoch_end(...)` method in the skeleton of the `F1ScoreModelCheckpoint` class for using the macro F1 score as the indicator for the best model. Save the best model in each epoch using the method `model.save`. Retrain the model using the parameters from 1.3 and **print** and **report** your new macro F1 score on the test set. (2P)

Hint: You may use `sklearn.metrics.f1_score(...)`.

1.5 Hyperparameter Optimization

(3P)

Experiment with the following hyperparameters: `hidden_units`, `dropout`, `batch_size`.

You may employ random hyperparameter optimization as usual. However, given the duration it takes to train a single configuration, it is fine to manually select a few promising hyperparameter configurations in this home exercise.

Report your three best hyperparameter sets (“best” according to the development set). Also **report** the macro F1 score of these three parameter sets on the test set.